

Contents

1	Introduction	3
1.1	Post Release Testing supports developers to efficiently do operations work	3
2	Continuous Delivery only covers practices until release	5
2.1	Continuous Delivery disregards security and operations topics . . .	5
2.2	Fast time to market is crucial	5
2.3	Continuous monitoring is hard	5
2.4	Simple day to day work must be automated	5
2.5	How to read this masterthesis	5
3	State of the art technologies and practices are the foundations for NPRT	6
3.1	The Continuous Delivery Pipeline consists of commit, automated testing and deployment	6
3.2	Docker packages applications	6
3.3	Kubernetes is a cluster operating system	6
3.4	Monitoring a highly dynamic infrastructure is role centric	6
4	Nonfunctional Production Regression Testing extends the Continuous Delivery Pipeline	7
5	Deployer implements and automates NPRT	14
5.1	Version centric testing via commit hashes	14
5.2	A canary and its testing metrics know about themselves	14
5.3	Controlling deploys in the pipeline and manually	14
5.4	Comparison of versions triggers webhooks for further actions . .	14
6	Evaluation	15
6.1	How NPRT changes the behaviour of development teams	15
6.1.1	Deploys	15
6.1.2	Cycletime	15

6.1.3	Change	15
6.1.4	True/False Positives/Negatives	15
6.2	NPRT compared to other in production testing strategies	15
6.2.1	Netflix Simian Army to intensify NPRT	15
6.2.2	Synthetic Monitoring is functional post release testing	15
7	Conclusion	16
7.1	Resume	16
7.2	Outlook and future work	16

Chapter 1

Introduction

1.1 Post Release Testing supports developers to efficiently do operations work

Many people talk about DevOps as well as there are multiple definitions and interpretations of the term DevOps. DevOps is referred as a philosophy, a culture, practices and specific tools. For my research, I will focus on two different aspects of the term DevOps:

The first one is the perspective of operation teams. Operation teams traditionally modeled infrastructure by installing physical hardware and by manually installing software components. With the rise of virtual machines and the cloud, it became possible to model infrastructure in software¹. Modelling via software enables operation teams to use tools and practices² as seen in software engineering. Infrastructure code is version controlled, tested and can be automatically deployed.

The other aspect of DevOps³ is the perspective of developer teams. Previously developer teams were only responsible for developing new features. Software engineering practices got established and proven. One of those practices is the continuous delivery pipeline⁴. The last step of the continuous delivery pipeline is the deployment. Formerly operation teams were responsible for deploying new features. The deployment as last step of the continuous delivery pipeline shifts a responsibility from operation to development. This shows that developer teams are becoming more and more responsible for running the software, they built.

¹“Infrastructure as Code” describes different dynamic infrastructure types [3, p. 30] and how to model those by code [3, p. 42].

²In the chapter “Software Engineering Practices for Infrastructure” [3, p. 179-194] practices like version controlling, continuous integration are described.

³The book “DevOps” [1] is written in the view of a developer running a system.

⁴For theoretical details on the continuous delivery pipeline read Part II of “Continuous Delivery” [2, p. 103-140] or a more practical approach by Wolff [4].

In the following the structure of the thesis is outlined. Every chapter is briefly discussed, what it is about.

In the first chapter we will walk through the foundations. The chapter mentions technologies, which are used for the thesis and gives references. Furthermore it gives references to the practices which are used and are crucial for the thesis. The references are properly selected, to understand the details if they are not known and to understand what the technologies and techniques are used for. In summary those are kubernetes, continuous delivery, continuous deployment and techniques from infrastructure as code and site reliability engineering.

The second chapter is conceptual macro view to the technique nonfunctional production regression testing. The text walks through the general environment and discusses the most important concepts and how they communicate with each other. The most important steps of the continuous delivery pipeline are discussed and it explains how the pipeline is extended in order to have the technique of non-functional production regression testing. The text argues how the methodologies of nonfunctional production are embedded in the pipeline and how the pipeline must be extended.

In the third chapter we will get to the concrete implementation of the non-functional production regression testing. The chapter will go into the details of how the concept is implemented. Concretely the software deployer is described, which was implemented in the context of this thesis.

Chapter four is about the evaluation of the new approach. We will investigate the use of the technique and customized software in two different companies. The first company is Gapfish, a four year old startup, and the software department of DIN, a company which is established for a hundred years. We are going to evaluate positive outcomes, still problematic concerns and their improvements. Another part of the evaluation is the comparison to other techniques which other companies and groups developed and tested. We differentiate in their features, advantages and disadvantages.

In the last chapter, the conclusion, the whole thesis is summarized and all the chapters are resumed. Important is the second part of the conclusion, in which we have an outlook to further improvements and how the technique can be extended to have further upgrades to delivery pipelines.

Chapter 2

Continuous Delivery only covers practices until release

2.1 Continuous Delivery disregards security and operations topics

2.2 Fast time to market is crucial

2.3 Continuous monitoring is hard

monitoring change and trying to predict the future from data

2.4 Simple day to day work must be automated

2.5 How to read this masterthesis

Chapter 3

State of the art technologies and practices are the foundations for NPRT

- 3.1 The Continuous Delivery Pipeline consists of commit, automated testing and deployment**
- 3.2 Docker packages applications**
- 3.3 Kubernetes is a cluster operating system**
- 3.4 Monitoring a highly dynamic infrastructure is role centric**

Chapter 4

Nonfunctional Production Regression Testing extends the Continuous Delivery Pipeline

Nonfunctional Production Regression Testing is the topic of the masterthesis and we will discuss the core of it in this chapter. The designation itself describes precisely what the technique and what the new practice is about. Therefore we will go shortly through every single term in the following.

Nonfunctional refers to the metrics, which we're evaluating in the test; These metrics are only nonfunctional and as a consequence generically applicable to multiple applications. The term production refers to the environment. The metrics, which are collected are collected in the production environment. And finally the term regression referring to the testing strategy. The metrics will be compared between two different versions and the latter version is tested for a regression, concretely a decline of the monitored metrics.

The testing technique provides some further features, which are not included in the designation. Indeed the testing technique is completely automatable and you can continuously apply it to the new versions. The testing technique is designed in respect to failing as fast as possible and inform developers.

When it comes to automation and continuous, the thoughtful reader is now probably reminded of continuous integration, delivery and deployment. This testing techniques evolved naturally from those practices and extends those. Those already established practices support the developers until the software is deployed. In contrast to that, nonfunctional production regression testing, supports the developers during after the deploy and while the software runs in production.

To understand the testing procedure in a whole and completely, it is necessary to show a complete overview of the whole testing and production environment. We will go through the steps of the pipeline and discuss it in a nutshell as you can

see in the figures.

The steps in the very beginning are known from the established practices continuous delivery and deployment respectively. But it is necessary to touch them and integrate them to the whole picture and outline the special characteristics for the new testing technique.

At the very beginning, there is a developer, who changed the code locally on his working machine. There is also a version control system. And in the first step the developer commits the code to the version control system. The main focus is on git. Subversion is possible, too though. In the following description we will stick to the terms and notions of git.

Next there is a continuous integration system. After the commit happened, the second step is a message to the continuous integration system. The message holds the information that a new commit exists and the continuous integration server clones the code from the version control system and checks out the specified version. Now the continuous integration system has three major jobs. The first one is to start a build process, the second is to run the tests and the third is to give the deploy signal.

In step four, namely running the build, it is typical to compile binaries, render assets and further artifacts. For our purposes it is especially necessary to build at least one or multiple docker images.

The import thing about this is, that we need to identify every docker image to a specific build. Therefore we use the commit hash, the version created by the version control system. This commit hash will follow us through the whole pipeline. This is important to be able to trace every step in the pipeline for a specific version. With this thought in mind, the docker image is tagged with the commit hash of this version and the name of the branch. Along the way it is mentioned, that the branch name is not absolutely necessary to definitely determine the version. The branch name is included for better readability for a developer and approximately recognize what the image version is about.

The continuous integration server then pushes the ready build docker image to an image registry, such as docker hub. Nevertheless this can be a private registry as well. This registry serves later as an artifact repository.

The second continuous integration step, or in total step six, the continuous integration server runs the tests. There can be multiple stages, such as unit, feature or smoke tests. Yet we do not need to recall all the details here.

The last step of the continuous integration system is to send a deploy signal. In the shown figure this is step seven. When you look at the tests, the result could on the one hand be a failure or on the other hand be successful. If the test have failed, the remaining pipeline will be cancelled and the developer will be informed. Just as you know it from a typical continuous integration system. If the tests have been successful and accordingly the build including all test stages have been successful,

the continuous integration system sends the signal to deploy.

It makes sense to deploy only specific versions and not every commit. The practice which is pretty common, is that you develop new features in a separate branch. For those version it is common to not send a deploy signal even though the branch build and tests are successful. Usually after there has been a review and a decision to deploy the changes to production, even though it is a very small change. But when the decision is made and merged into a specified branch, for instance the master branch, this version will go to production.

However, just to clarify, each built image for every single version is sent, independently of successful tests and independently of the intention to go to production, to the image repository. The reason could be a staging system and even running the tests inside the build image. But this is just a side note.

So the deploy signal is given when two requirements are fulfilled: the build and tests are successful and it is a version which is planned to go to production.

Until this point, as it was already mentioned, it is just a usual continuous delivery or deployment pipeline, which is commonly used in the development process. But since nonfunctional production regression testing is a technique, which is supposed to be completely automated, such a prior described delivery pipeline including automated deployments is precondition. From now on it is becoming interesting, hence the testing technique supports the developers post deployment in production instead of the old practices before the deploy.

The next unit is the deployer. It is the software, which is particularly implemented for this masterthesis. In the next chapter the deployer is described in detail. This chapter demonstrates how the deployer is embedded in the pipeline or in other words in the environment amongst all other tools. In the meantime tools exist, which have a similar purpose. It is crucial to have full control over the whole deployment process and as a consequence it was necessary to implement the software and have it customizable.

We could also implement the logic of the deploy deploy in the continuous integration system. But we had to decide against that, because the deploy needs full access to the production system and the continuous integration system is in our case outsourced to a third party company. We don not want to give other companies full access to another company. However this meant, that we had to implement some steps again, which a continuous integration server already implements.

The deploy message, which the continuous integration system sends to the deployer, includes the commit hash again. And again we use the commit hash to identify the version. Now the deployer executes three major steps: In the first step the deployer fetches the code from the version control system. This is the same thing, which the continuous integration system does. The repository is necessary, because it holds the files which describe our infrastructure. We want to version

control the definitions of our infrastructure in order to be able to relate the version of the infrastructure to the version of the code and the version of the artifacts. At kubernetes those infrastructure definitions are made up of different resources, which were already mentioned in the foundations chapter.

In the second step deployer modifies those infrastructure resources. The infrastructure should use the related docker image and be aware of its own version.

And the third step or the ninth step in the figure, deployer applies the modified resources to the production system, which is shown in the figure as well. We note that not only application code changes are deployed, but also infrastructure changes. We note that we deploy the infrastructure changes in a continuous fashion as well.

Next we go on with the process of how the production system updates itself. The production system is a kubernetes cluster, which the previous chapter foundations describes. For now, we assume, that the production system runs a typical three tier webapp. The webapp is made up out of a loadbalancer, multiple stateless web servers and a stateful database. The loadbalancer balances the requests between the web servers. And the web servers communicate with the database. And the database consists of a replicated cluster.

We explain firmly how the typical three tier webapp is implemented in kubernetes. For more details refer to the references given in the foundations chapter. With kubernetes we define the mentioned elements with the kubernetes resources: service, deployment and its pods and statefulset. For simplification, we imagine the service as a loadbalancer. Then there is a deployment, which manages the existence of the identical pods. And lastly the statefulset, which manages the stateful pods, with their unique name and disk.

So the loadbalancer receives a request from the client. The loadbalancer selects a pod via round robin and proxies the request to the pod. The pod probably communicates with the database and sends the request back to the client, where the loadbalancer acts again as proxy.

We now look at the change, which happens to the production system. Earlier, deployer made changes to the definitions of the kubernetes resources in the deployers memory and communicates those changes to the production system. We are interested in particular in the deployments. As an illustration we only look at what happens to the deployment and the pods.

The deployer talks to the kubernetes master api and sends the changes. The master manages the concrete changes. It swaps out one pod by another by stopping the pods in the old version and starting the pods in the new version. The procedure is called rolling update. Now the pods are all swapped out and run in the new version.

Another part of the cluster are the monitoring agents. They pick up monitoring metrics in different ways. They collect the data, we are interested in for

the regression test in production, as well. We are talking about the nonfunctional metrics. We selected the metrics, defined by the four golden metrics of google's sre (identified). The metrics are throughput, latency, errorrate and utilization.

With utilization we are lucky, because kubernetes already implements a collection of cpu and memory. But we need to instrument the application to collect the other metrics, throughput, latency and errorrate. So the pods send the instrumentation data to the monitoring agent. The monitoring agent is an statsd server, which collection the data and aggregates the data and forwards it to the monitoring system.

The monitoring system consists basically out of a timeseriesdatabase, a graphing user interface and a alarm system. The timeseriesdatabase persists the metrics. And the user can define graphs from those metrics, which the user interface presents visually. You can define rules in the alarm system, which monitor the metrics in the timeseriesdatabase and then, in case the rule is violated, sends notifications.

We are interested in the monitoring data of specific versions. Consequently the monitoring agents needs to send the monitoring labeled with the specific version to the monitoring system. This is important, since we want to compare the metrics of the different version.

We have different possibilities to compare those version. One possibility is, that we compare current and historical data. For instance to compare the metrics of the current production system with the metrics of the production system of the day before or even the week before and compare the different versions of those times.

We are following a different approach, because when we are comparing the current production system with the production system of last week, we have lots of different changes. The current traffic must not be the same traffic as last week, the load of the production system must not be the same load and other system with which the application is interacting with must not be the same.

That is why we decided to compare two different versions which run in the production system concurrently. This brings not only the advantage, that you have the very same traffic, but also the advantage, that there is less risk involved. We illustrate the advantage of less risk now by demonstrating the process of deploying the second version and comparing it to the old version.

Ok, if you compare the two versions with each other, you will do at as follows. Deployer create another deployment resource from the one that already exists. Deployer calls this other deployment resource canary deployment. The creation of the canary deployment resource has the effect, that not only pods of version I are running in production, but there is pods running in version II as well. Similar to the regular deployment, the canary deployment defines how many pods in which version are supposed to be running.

We want to test, if there is a regression respectively a degradation between the two versions. On account to the fact, that a regression is possible and when introducing change, a regression is very likely, we at least want to affect as little users. So what do we do for that? In our example there three pods running in version I and only one pod in version II. This is a ratio of three to one and due to the fact that the loadbalancer uses round robin as the scheduling algorithm, only one in four requests, so 25% of the total traffic is sent to the pods in version II, which is to test.

This certainly lowers the risk of failure and that users are affected by a regression. Even if the request of specific single user hits the degraded pod, the next request of the same user has the probability 75% to hit the old stable version.

A limitation to this technique is that the new version II needs to be able to run side by side with the old version I. In most cases, that means that the new version needs to be semantically almost identical to the old version. So version II should not provide functional changes compared to version I, but only nonfunctional changes. However that means we cannot test new feature like in an A/B test. Instead we can test performance improvements, refactoring or updates.

They call this technique canary releasing. Again, you change would only change a part of the production system, the canary instead of the whole. Devops TODO examines this technique in more detail.

Assuming we would want to test features in production, the current implementation of the technique is not suitable. If we wanted to do that, we would need to include the loadbalancer. The loadbalancer would need to remember which user is proxied to which version, so that the next request of that user goes to the same version, thus the user sees the same set of features as before. The design of the database could potentially be affected as well and could be needed to be loadbalanced for the users. The technique we just described is usually called an A/B test. The disadvantage of the A/B test is that the same user will hit on the same potentially degraded service and it is not that simple to automatically provide a stable service to the user. Due to simplification, we did decide to not include the implementation of the loadbalancing.

We want to state that it is suboptimal to run multiple versions in the cluster like also mentioned in devops TODO. Rolling updates require it to be able to have two versions in production, though. And kubernetes utilizes rolling updates as a technique to provide zero downtime deployments. Accordingly our proposed technique does not introduce a worsening to that. But as in devops mentioned, you should avoid to run more than two versions at the same time in production. Deployer ensures that by either updating a deployment, creating a canary deployment, or creating a deployment in a new version, just before it deleted the canary deployment.

Especially to test the latter, security updates, is absolutely appealing, since we

can fully automate the procedure of updating the dependencies of our application in a fully automated and in a way, which would have a very low risk. We could have a job, which checks frequently for any new version, pushes the updates to the version control system, the continuous integration system runs the pre deploy tests, deployer deploys the update and even in production we check the update for an regression. We could save a lot of developer time, who would usually need to take care of the whole updating procedure. And even if there is a degradation in production, a small amount requests is affected, because we send only a reasonable amount of traffic, which arrives at the same time, to the potentially degraded version. Further more we limit the time the degraded version is in production, because we automate detection of the degradation and the rollback to the old stable version.

Now version I and version II send metrics via the monitoring agent to the monitoring system. We tag the metrics with the specific version as well. The monitoring system stores the two comparable metrics of the two different version in the timeseriesdatabase. Now you can define a graph to the metrics of version I and you can too define a graph of metrics in version II. We then compare the two graphs by for instance subtracting one from the other and monitor the result. We let this running for a specified time in production. We need to decide on how long we want to compare the versions. That depends on how much traffic is in production, because when we would few traffic in production, we wanted to compare for a longer time. We suggest to have a well balanced test scenario in terms of load.

We do not need to generate the test traffic, we do not need to weight traffic and we do not need to think about edge cases. These are all advantages, that we get for free from the production traffic. We save time and work, because the users generate the test data, instead of us. The users create more requests and with that test data for parts of the application, which are more important. Consequently the users reasonably weight the test data. And lastly the longer we run the comparison in production, users will produce more of those edge cases, which would be hard to make up.

We are aware of that the two compared versions do not receive the very same requests. Hence the comparison is not perfect. In future work we could extend the technique to achieve that. We could simply clone the requests, send the original request to the stable version and send a cloned request to the canary. The load-balancer could then differentiate between the two responses of the two versions. We would reject the response of the canary. And we would forward the response of the stable version.

As a result we even lower the risk, because the potentially degraded version does not even respond to real users. Ergo we do not have any risk of a degradation of our production service which we cause by testing the new version.

Chapter 5

Deployer implements and automates NPRT

5.1 Version centric testing via commit hashes

build, test, deploy, only 2 versions in production. undeploy a canary.

5.2 A canary and its testing metrics know about themselves

5.3 Controlling deploys in the pipeline and manually

5.4 Comparison of versions triggers webhooks for further actions

monitoring validation, fail

Chapter 6

Evaluation

6.1 How NPRT changes the behaviour of development teams

6.1.1 Deploys

6.1.2 Cycletime

6.1.3 Change

6.1.4 True/False Positives/Negatives

6.2 NPRT compared to other in production testing strategies

6.2.1 Netflix Simian Army to intensify NPRT

6.2.2 Synthetic Monitoring is functional post release testing

Chapter 7

Conclusion

7.1 Resume

7.2 Outlook and future work

Bibliography

- [1] Len Bass, Ingo Weber, and Liming Zhu. *DevOps: a software architect's perspective*. Addison-Wesley, 2015.
- [2] Jez Humble and David Farley. *Continuous Delivery: reliable software releases through build, test, and deployment automation*. Addison-Wesley, 2010.
- [3] Kief Morris. *Infrastructure as code : managing servers in the cloud*. O'Reilly, 2016.
- [4] Eberhard Wolff. *Continuous Delivery: der pragmatische Einstieg*. 2. dpunkt.verlag, 2016.